

## ✓ Appendix 1 - Preliminary analysis

```
!pip install pandas numpy seaborn matplotlib
```

### Load and Clean the Data

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from google.colab import files

# Load the dataset
df = pd.read_csv("cstads_2122_pumf.csv")

# Dataset overview
print("Original DataFrame:\n")
print(df.shape)
df.describe()
df.info()

# Check the first few rows
print("\n", df.head())

# Remove duplicates if any
df = df.drop_duplicates()

# Standardize categorical values (example: converting 'Male'/'M' variations)
df['DVGENDER'] = df['DVGENDER'].replace({1: 'Female', 2: 'Male'})
df['DVURBAN'] = df['DVURBAN'].replace({1: 'Urban', 2: 'Rural'})
df['PROVID'] = df['PROVID'].replace({10: 'NL', 11: 'PEI', 12: 'NS', 24: 'QC', 35: 'ON', 46: 'MN', 47: 'SK', 48: 'AB', 59: 'BC'})

# Identify missing values
df.replace({96: pd.NA, 98: pd.NA, 99: pd.NA, 996: pd.NA, 999: pd.NA}, inplace=True)
print("\nMissing Values:\n", df.isnull().sum())

# Calculate the percentage of missing values per column
missing_percentage = df.isnull().mean() * 100

# Drop columns with more than 75% missing values
df_cleaned = df.loc[:, missing_percentage <= 75]

# Cleaned dataset overview
print("\nCleaned DataFrame (columns with more than 75% missing values removed):\n")
print(df_cleaned.shape)
df_cleaned.describe()
df_cleaned.info()

# Check the first few rows
print("\n", df_cleaned.head())

# Identify missing values
print("\nMissing Values:\n", df_cleaned.isnull().sum())

# Move cleaned dataframe to df for convenience
df = df_cleaned
```

↔ Original DataFrame:

```
(61096, 168)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 61096 entries, 0 to 61095
Columns: 168 entries, SEQID to DVAVCIGD
dtypes: float64(1), int64(167)
memory usage: 78.3 MB
```

	SEQID	PROVID	GRADE	DVGENDER	DVURBAN	DVRES	DVORIENT	DVDESCRIBE	\
0	18338	11	9	2	1	3	2	1	
1	16111	24	10	1	1	1	2	4	
2	20587	10	7	2	1	1	2	1	
3	54568	12	8	99	2	1	1	1	
4	40991	10	7	1	2	1	2	1	

	WTPUMF	GH_010	...	BUL_100	BUL_110	BUL_120	DVTY1ST	DVTY2ST	\
0	1.164716	2	...	2	2	4	3	7	
1	36.141169	3	...	2	2	1	3	7	
2	1.448578	2	...	2	2	1	3	7	
3	3.036660	4	...	2	2	2	3	7	
4	3.745233	2	...	99	99	99	3	7	

	DVLAST30	DVAMTSMK	DVCIGWK	DVNDSMK	DVAVCIGD
0	2	96	996	96	96
1	2	96	996	96	96
2	2	96	996	96	96
3	2	96	996	96	96
4	2	96	996	96	96

[5 rows x 168 columns]

Missing Values:  
SEQID 5  
PROVID 0  
GRADE 0  
DVGENDER 5025  
DVURBAN 0  
...  
DVLAST30 375  
DVAMTSMK 57291  
DVCIGWK 56879  
DVNDSMK 56876  
DVAVCIGD 56876  
Length: 168, dtype: int64

Cleaned DataFrame (columns with more than 75% missing values removed):

(61096, 138)  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 61096 entries, 0 to 61095  
Columns: 138 entries, SEQID to DVLAST30  
dtypes: float64(1), int64(1), object(136)  
memory usage: 64.3+ MB

	SEQID	PROVID	GRADE	DVGENDER	DVURBAN	DVRES	DVORIENT	DVDESCRIBE	WTPUMF	\
0	18338	PEI	9	Male	Urban	3	2	1	1.164716	
1	16111	QC	10	Female	Urban	1	2	4	36.141169	

Generate Summary Statistics

```
# Summary statistics for numerical variables
print("Summary Statistics:\n", df.describe())

# Count unique values for categorical variables
print("\nCategorical Data Distribution:\n", df[['DVGENDER']].apply(pd.Series.value_counts))
print("\n", df[['PROVID']].apply(pd.Series.value_counts))
print("\n", df[['GRADE']].apply(pd.Series.value_counts))
print("\n", df[['DVURBAN']].apply(pd.Series.value_counts))
```

Summary Statistics:

	GRADE	WTPUMF
count	61096.000000	61096.000000
mean	9.141548	35.353820
std	1.622969	56.903468
min	7.000000	0.507389
25%	8.000000	6.075865
50%	9.000000	25.260619
75%	10.000000	38.635371
max	12.000000	1561.260211

Categorical Data Distribution:

DVGENDER

Male	28903
Female	27168

PROVID

QC	10863
AB	9260
NL	7032
NS	6999
BC	6885
ON	6745
SK	5596
PEI	4616
MN	3100

GRADE	
8	12500
7	12436
9	11055
10	10218
11	8859
12	6028

DVURBAN	
Urban	49577
Rural	11519

Perform Correlation Analysis

```
import re

# Define the prefixes to exclude
exclude_prefixes = ["UND", "MET", "XTC", "HAL", "HER", "COC", "SYN", "BZP", "TNB", "TRP", "GLU", "SAL",
                    "SLP", "STI", "DEX", "GRV", "SED", "POLY", "PH", "DR", "BEH", "BUL"]

# Create a regex pattern to match column names starting with these prefixes
pattern = re.compile(r'^(?:' + '|'.join(exclude_prefixes) + r').*')


# Identify columns to drop
columns_to_drop = [col for col in df.columns if pattern.match(col)]

# Drop the unwanted columns
df_filtered = df.drop(columns=columns_to_drop)

# Print the removed columns
print("Removed columns:", columns_to_drop)

# Remaining dataset overview
print("\nCleaned DataFrame (columns not associated with smoking/alcohol/cannabis removed):\n")
print(df_filtered.shape)
df_filtered.describe()
df_filtered.info()

# Move cleaned dataframe to df for convenience
df = df_filtered
```



Removed columns: ['UND\_010', 'UND\_020', 'MET\_010', 'XTC\_010', 'HAL\_010', 'HER\_010', 'COC\_010', 'SYN\_010', 'BZP\_010', 'TNB\_010', 'TRP\_010', 'GLU\_010', 'SAL\_010', 'SLP\_010', 'STI\_010', 'DEX\_010', 'GRV\_010', 'SED\_010', 'POLY\_010', 'PH\_010', 'DR\_010', 'BEH\_010', 'BUL\_010']

Cleaned DataFrame (columns not associated with smoking/alcohol/cannabis removed):

(61096, 58)

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 61096 entries, 0 to 61095

Data columns (total 58 columns):

#	Column	Non-Null Count	Dtype
0	SEQID	61091 non-null	object
1	PROVID	61096 non-null	object
2	GRADE	61096 non-null	int64
3	DVGENDER	56071 non-null	object
4	DVURBAN	61096 non-null	object
5	DVRES	59793 non-null	object
6	DVORIENT	50683 non-null	object
7	DVDESCRIBE	58548 non-null	object
8	WTPUMF	61096 non-null	float64
9	GH_010	60248 non-null	object
10	GH_020	60160 non-null	object
11	SS_010	61011 non-null	object
12	TS_011	60834 non-null	object
13	TP_016	59237 non-null	object
14	TP_046	58940 non-null	object
15	TP_056	59433 non-null	object
16	TP_066	59263 non-null	object
17	TP_086	59242 non-null	object
18	ELC_026a	59670 non-null	object
19	ELC_026b	58644 non-null	object
20	ELC_026c	58606 non-null	object
21	VAP_010	60487 non-null	object
22	CI_010	59877 non-null	object
23	VAP_020	59557 non-null	object
24	VAP_030	58820 non-null	object
25	VAP_040	58388 non-null	object

26	VAP_050a	59872	non-null	object
27	VAP_050b	58900	non-null	object
28	VAP_060	59607	non-null	object
29	ALC_010	60525	non-null	object
30	ALC_020	26907	non-null	object
31	ALC_030	27582	non-null	object
32	ALC_040	26722	non-null	object
33	ALC_050	27874	non-null	object
34	NRG_010	60047	non-null	object
35	NRG_020	59039	non-null	object
36	NRG_030	58854	non-null	object
37	NRG_040	58833	non-null	object
38	NRG_050	58906	non-null	object
39	ALC_075	25943	non-null	object
40	CAN_010	60573	non-null	object
41	CAN_130	59884	non-null	object
42	CAN_140	59272	non-null	object
43	BS_010	59810	non-null	object
44	PR_100	60001	non-null	object
45	PR_030	59984	non-null	object
46	PR_050	59950	non-null	object

```
# Compute the correlation matrix
# Convert columns to numeric, errors='coerce' will handle non-numeric values
correlation_matrix = df.apply(pd.to_numeric, errors='coerce').corr()

# Heatmap visualization
plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix, annot=True, cmap="coolwarm", fmt=".2f", linewidths=0.5)
plt.title("Feature Correlation Heatmap")
plt.show()

# Pairplot for key variables
# Convert columns to numeric for pairplot as well
sns.pairplot(df[['GRADE', 'SS_010', 'ALC_010', 'CAN_010']].apply(pd.to_numeric, errors='coerce'))
plt.show()
```



Feature Correlation Heatmap

