

## ✓ Appendix 1 - Final Results and Code

```
!pip install pandas numpy seaborn matplotlib
```

 [Show hidden output](#)

```
pip install ydata-profiling
```

```
 Requirement already satisfied: typeguard<5,>=3 in /usr/local/lib/python3.11/dist-packages (from ydata-profiling) (4.4.2)
Collecting imagehash==4.3.1 (from ydata-profiling)
  Downloading ImageHash-4.3.1-py2.py3-none-any.whl.metadata (8.0 kB)
Requirement already satisfied: wordcloud>=1.9.3 in /usr/local/lib/python3.11/dist-packages (from ydata-profiling) (1.9.4)
Collecting dacite>=1.8 (from ydata-profiling)
  Downloading dacite-1.9.2-py3-none-any.whl.metadata (17 kB)
Requirement already satisfied: numba<=0.61,>=0.56.0 in /usr/local/lib/python3.11/dist-packages (from ydata-profiling) (0.60.0)
Collecting PyWavelets (from imagehash==4.3.1->ydata-profiling)
  Downloading pywavelets-1.8.0-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (9.0 kB)
Requirement already satisfied: pillow in /usr/local/lib/python3.11/dist-packages (from imagehash==4.3.1->ydata-profiling) (11.1.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.11/dist-packages (from jinja2<3.2,>=2.11.1->ydata-profiling)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib<=3.10,>=3.5->ydata-profiling)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.11/dist-packages (from matplotlib<=3.10,>=3.5->ydata-profiling)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib<=3.10,>=3.5->ydata-profiling)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib<=3.10,>=3.5->ydata-profiling)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib<=3.10,>=3.5->ydata-profiling)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib<=3.10,>=3.5->ydata-profiling)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.11/dist-packages (from matplotlib<=3.10,>=3.5->ydata-profiling)
Requirement already satisfied: llvmlite<0.44,>=0.43.0dev0 in /usr/local/lib/python3.11/dist-packages (from numba<=0.61,>=0.56.0->ydata-profiling)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas!=1.4.0,<3.0,>1.1->ydata-profiling)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas!=1.4.0,<3.0,>1.1->ydata-profiling)
Requirement already satisfied: joblib>=0.14.1 in /usr/local/lib/python3.11/dist-packages (from phik<0.13,>=0.11.1->ydata-profiling)
Requirement already satisfied: annotated-types>=0.6.0 in /usr/local/lib/python3.11/dist-packages (from pydantic<=2->ydata-profiling)
Requirement already satisfied: pydantic-core==2.33.1 in /usr/local/lib/python3.11/dist-packages (from pydantic<=2->ydata-profiling)
Requirement already satisfied: typing-extensions>=4.12.2 in /usr/local/lib/python3.11/dist-packages (from pydantic<=2->ydata-profiling)
Requirement already satisfied: typing-inspection>=0.4.0 in /usr/local/lib/python3.11/dist-packages (from pydantic<=2->ydata-profiling)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.24.0->ydata-profiling)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.24.0->ydata-profiling)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.24.0->ydata-profiling)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.24.0->ydata-profiling)
Requirement already satisfied: patsy>=0.5.6 in /usr/local/lib/python3.11/dist-packages (from statsmodels<1,>=0.13.2->ydata-profiling)
Requirement already satisfied: attrs>=19.3.0 in /usr/local/lib/python3.11/dist-packages (from visions<0.8.2,>=0.7.5->visions[type_image_path]<0.8.2,>=0.7.5->ydata-profiling)
Requirement already satisfied: networkx>=2.4 in /usr/local/lib/python3.11/dist-packages (from visions<0.8.2,>=0.7.5->visions[type_image_path]<0.8.2,>=0.7.5->ydata-profiling)
Collecting puremagic (from visions<0.8.2,>=0.7.5->visions[type_image_path]<0.8.2,>=0.7.5->ydata-profiling)
  Downloading puremagic-1.28-py3-none-any.whl.metadata (5.8 kB)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.7->matplotlib<=3.10,>=3.5)
Downloading ydata_profiling-4.16.1-py2.py3-none-any.whl (400 kB)
 400.1/400.1 kB 12.5 MB/s eta 0:00:00
Downloading ImageHash-4.3.1-py2.py3-none-any.whl (296 kB)
 296.5/296.5 kB 20.3 MB/s eta 0:00:00
Downloading dacite-1.9.2-py3-none-any.whl (16 kB)
Downloading multimethod-1.12-py3-none-any.whl (10 kB)
Downloading phik-0.12.4-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (687 kB)
 687.8/687.8 kB 32.2 MB/s eta 0:00:00
Downloading visions-0.8.1-py3-none-any.whl (105 kB)
 105.4/105.4 kB 8.8 MB/s eta 0:00:00
Downloading puremagic-1.28-py3-none-any.whl (43 kB)
 43.2/43.2 kB 3.4 MB/s eta 0:00:00
Downloading pywavelets-1.8.0-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (4.5 MB)
 4.5/4.5 MB 47.7 MB/s eta 0:00:00
Building wheels for collected packages: htmlmin
  Building wheel for htmlmin (setup.py) ... done
  Created wheel for htmlmin: filename=htmlmin-0.1.12-py3-none-any.whl size=27081 sha256=3020ad49a2290d50f616ceefb6a1440db14ae1f5d0cfe
  Stored in directory: /root/.cache/pip/wheels/8d/55/1a/19cd535375ed1ede0c996405ebffe34b196d78e2d9545723a2
Successfully built htmlmin
Installing collected packages: puremagic, htmlmin, PyWavelets, multimethod, dacite, imagehash, visions, phik, ydata-profiling
Successfully installed PyWavelets-1.8.0 dacite-1.9.2 htmlmin-0.1.12 imagehash-4.3.1 multimethod-1.12 phik-0.12.4 puremagic-1.28 visio
```

### Load and Review the Data

```
import pandas as pd
import numpy as np
from google.colab import files
```

```
# Load the dataset
df = pd.read_csv("cstads_2122_pumf.csv")
```

```
# Dataset overview
```

```
print("Original DataFrame:\n")
print(df.shape)
df.describe()
df.info()
```

Original DataFrame:

```
(61096, 168)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 61096 entries, 0 to 61095
Columns: 168 entries, SEQID to DVAVCIGD
dtypes: float64(1), int64(167)
memory usage: 78.3 MB
```

**Primary Dataset Preparation** Data pre-processing including: identifying and handling missing or duplicate values, feature selection, creating a balanced dataset, graphically visualizing data, identifying correlations, and explaining them.

```
# Remove duplicates if any
df = df.drop_duplicates()

# Replace survey-specific missing value codes with pd.NA
df.replace({96: pd.NA, 98: pd.NA, 99: pd.NA, 996: pd.NA, 999: pd.NA}, inplace=True)

# Calculate percentage of missing values per column
missing_percentage = df.isnull().mean() * 100

# Drop columns with more than 50% missing values
df_cleaned = df.loc[:, missing_percentage <= 50]

# Move cleaned dataframe to df
df = df_cleaned

# Drop rows with any remaining missing values
df.dropna(inplace=True)

# Output updated dataframe info
print("Updated dataframe shape:\n", df.shape)
print("Updated dataframe features:\n", df.dtypes)
```

Updated dataframe shape:  
(23654, 133)  
Updated dataframe features:

SEQID	object
PROVID	int64
GRADE	int64
DVGENDER	object
DVURBAN	int64
...	...
BUL_110	object
BUL_120	object
DVTY1ST	object
DVTY2ST	object
DVLAST30	object
Length: 133, dtype: object	

```
import re
```

```
# Define the prefixes to exclude
exclude_prefixes = ["UND", "MET", "XTC", "HAL", "HER", "COC", "SYN", "BZP", "TNB", "TRP", "GLU", "SAL",
                    "SLP", "STI", "DEX", "GRV", "SED", "POLY", "DR", "BUL"]

# Create a regex pattern to match column names starting with these prefixes
pattern = re.compile(r'^(?:' + '|'.join(exclude_prefixes) + r').*')

# Identify columns to drop
columns_to_drop = [col for col in df.columns if pattern.match(col)]

# Drop the unwanted columns
df_filtered = df.drop(columns=columns_to_drop)

# Print the removed columns
print("Removed columns:", columns_to_drop)

# Remaining dataset overview
print("\nCleaned DataFrame (columns not associated with smoking/alcohol/cannabis removed):\n")
print(df_filtered.shape)
```

```
df_filtered.describe()
df_filtered.info()

# Move cleaned dataframe to df for convenience
df = df_filtered
```

```
17 TP_086      23654 non-null object
18 ELC_026a    23654 non-null object
19 ELC_026b    23654 non-null object
20 ELC_026c    23654 non-null object
21 VAP_010     23654 non-null object
22 CI_010      23654 non-null object
23 VAP_020     23654 non-null object
24 VAP_030     23654 non-null object
25 VAP_040     23654 non-null object
26 VAP_050a    23654 non-null object
27 VAP_050b    23654 non-null object
28 VAP_060     23654 non-null object
29 ALC_010     23654 non-null object
30 NRG_010     23654 non-null object
31 NRG_020     23654 non-null object
32 NRG_030     23654 non-null object
33 NRG_040     23654 non-null object
34 NRG_050     23654 non-null object
35 CAN_010     23654 non-null object
36 CAN_130     23654 non-null object
37 CAN_140     23654 non-null object
38 BS_010      23654 non-null object
39 PR_100      23654 non-null object
40 PR_030      23654 non-null object
41 PR_050      23654 non-null object
42 PR_060      23654 non-null object
43 PR_110      23654 non-null object
44 PH_010      23654 non-null object
45 PH_020      23654 non-null object
46 PH_030      23654 non-null object
47 PH_040      23654 non-null object
48 PH_051      23654 non-null object
49 PH_061      23654 non-null object
50 PH_052      23654 non-null object
51 PH_062      23654 non-null object
52 PH_110      23654 non-null object
53 PH_120      23654 non-null object
54 PH_070      23654 non-null object
55 PH_080      23654 non-null object
56 PH_130      23654 non-null object
57 PH_140      23654 non-null object
58 PH_090      23654 non-null object
59 PH_100      23654 non-null object
60 CA_020      23654 non-null object
61 ELC_041     23654 non-null object
62 ELC_042     23654 non-null object
63 ALC_080     23654 non-null object
64 CAN_050     23654 non-null object
65 PR_090      23654 non-null object
66 BEH_010     23654 non-null object
67 BEH_020     23654 non-null object
68 BEH_030     23654 non-null object
69 BEH_040     23654 non-null object
70 DVTY1ST     23654 non-null object
71 DVTY2ST     23654 non-null object
72 DVLAST30    23654 non-null object
dtypes: float64(1), int64(3), object(69)
memory usage: 13.4+ MB
```

## Exploratory Data Analysis

### Summary Statistics

```
# Count unique values for categorical variables
print("\nCategorical Data Distribution:\n", df[['DVGENDER']].apply(pd.Series.value_counts))
print("\n", df[['PROVID']].apply(pd.Series.value_counts))
print("\n", df[['GRADE']].apply(pd.Series.value_counts))
print("\n", df[['DVURBAN']].apply(pd.Series.value_counts))
```

```
Categorical Data Distribution:
DVGENDER
DVGENDER
2      12464
1      11190
```

	PROVID
PROVID	
24	4197
48	3904
10	2879
12	2648
59	2425
47	2287
35	2266
11	1909
46	1139

	GRADE
GRADE	
8	4434
9	4347
10	4163
11	3956
7	3893
12	2861

	DVURBAN
DVURBAN	
1	19418
2	4236

## Explore Distributions and Relationships

```
import seaborn as sns
import matplotlib.pyplot as plt
```


```
# Explore distributions
ax = sns.countplot(x='SS_010', data=df)
ax.set_xticklabels(['Yes', 'No'], rotation=0)
plt.xlabel('Ever tried cigarettes')
plt.ylabel('Count')
plt.show()
```

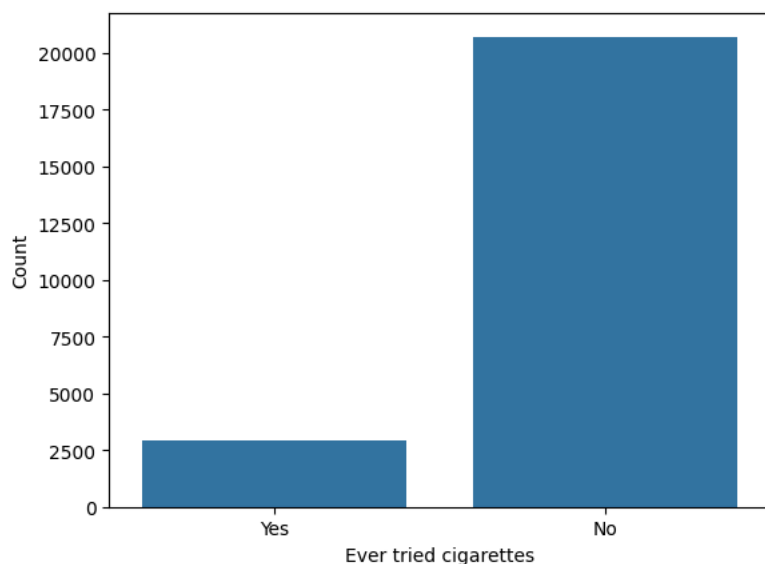
```
# Explore distributions
ax = sns.countplot(x='CI_010', data=df)
ax.set_xticklabels(['Never tried either', \
                    'Only tried cig(s)', \
                    'Only tried e-cig(s)', \
                    'Tried both / cig first', \
                    'Tried both / e-cig first', \
                    'Do not recall'], rotation=70)
plt.xlabel('Tried cigarettes or vapes first')
plt.ylabel('Count')
plt.show()
```

```
# Explore distributions
ax = sns.countplot(x='CA_020', data=df)
ax.set_xticklabels(['Very difficult', 'Fairly difficult', \
                    'Fairly easy', 'Very easy', 'I do not know'], rotation=70)
plt.xlabel('Difficulty getting cigarettes')
plt.ylabel('Count')
plt.show()
```


```
# explore distributions
ax = sns.countplot(x='DVTY2ST', data=df)
ax.set_xticklabels(['Current Daily Smoker', 'Current Occasional Smoker', \
                    'Former Smoker', 'Experimental Smoker', \
                    'Past Experimental Smoker', 'Puffer', 'Never Tried'], rotation=70)
plt.xlabel('Detailed smoking classification')
plt.ylabel('Count')
plt.show()
```

```
# explore distributions
ax = sns.countplot(x='TS_011', data=df)
ax.set_xticklabels(['Definitely yes (1)', 'Probably yes (2)', 'Probably not (3)', 'Definitely not (4)'], rotation=20)
plt.xlabel('Predicted use of cigarettes')
plt.ylabel('Count')
plt.show()
```


 <ipython-input-233-0b987d4f3c19>:6: UserWarning: set\_ticklabels() should only be used with a fixed number of ticks, i.e. after set\_t  
ax.set\_xticklabels(['Yes', 'No'], rotation=0)




<ipython-input-233-0b987d4f3c19>:13: UserWarning: set\_ticklabels() should only be used with a fixed number of ticks, i.e. after set\_t  
ax.set\_xticklabels(['Never tried either', \




<ipython-input-233-0b987d4f3c19>:25: UserWarning: set\_ticklabels() should only be used with a fixed number of ticks, i.e. after set\_t  
ax.set\_xticklabels(['Very difficult', 'Fairly difficult', \






```
<ipython-input-233-0b987d4f3c19>:33: UserWarning: set_ticklabels() should only be used with a fixed number of ticks, i.e. after set_t  
ax.set_xticklabels(['Current Daily Smoker', 'Current Occasional Smoker', \
```



```
<ipython-input-233-0b987d4f3c19>:42: UserWarning: set_ticklabels() should only be used with a fixed number of ticks, i.e. after set_t  
ax.set_xticklabels(['Definitely yes (1)', 'Probably yes (2)', 'Probably not (3)', 'Definitely not (4)'], rotation=20)
```

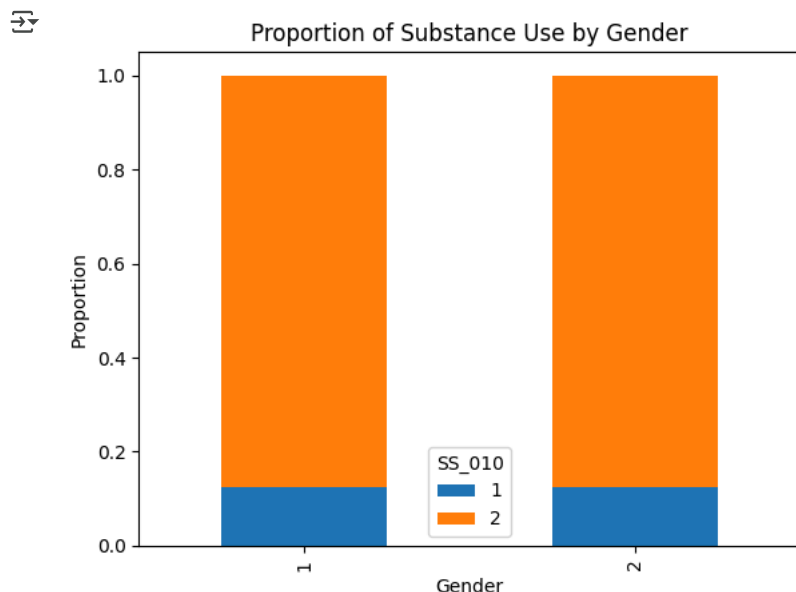


```
# Create a cross-tabulation (contingency table)
ct = pd.crosstab(df['DVGENDER'], df['SS_010'], normalize='index')

# Plot a stacked bar chart
ct.plot(kind='bar', stacked=True)

# Add labels and title
plt.xlabel('Gender')
plt.ylabel('Proportion')
plt.title('Proportion of Substance Use by Gender')

# Show the plot
plt.show()
```



## Generate EDA Report

```
from ydata_profiling import ProfileReport

# Create the profile
profile = ProfileReport(df, title="CSTADS EDA Report", explorative=True)

# Save to an HTML file
profile.to_file("cstads_eda_report.html")
files.download("cstads_eda_report.html")
```

 [Show hidden output](#)

## Adjust Data Types

Most machine learning models cannot directly handle categorical data stored as object dtype in pandas. They typically require numerical input. Therefore, categorical data must be encoded before being used in models. However, in CSTADS, categorical data is stored as object but already contains numbers that represent meaningful categories, as such we will convert the columns back to numeric format without altering their values.

```
# Identify object columns
object_cols = df.select_dtypes(include=['object']).columns
print("Object columns before conversion:", object_cols.tolist())

# Convert object columns to numeric only if they contain numeric values
for col in object_cols:
    # Drop NaNs, convert to string, and check if all non-null values are numeric
    if df[col].dropna().astype(str).str.isnumeric().all():
        df[col] = pd.to_numeric(df[col])

# Verify dtypes after conversion
```

```
print("\nData types after conversion:")
print(df.dtypes)
```

➦ Object columns before conversion: ['SEQID', 'DVGENDER', 'DVRES', 'DVORIENT', 'DVDESCRIBE', 'GH\_010', 'GH\_020', 'SS\_010', 'TS\_011', 'TP\_011']

```
Data types after conversion:
SEQID      int64
PROVID     int64
GRADE      int64
DVGENDER   int64
DVURBAN     int64
...
BEH_030    int64
BEH_040    int64
DVTY1ST    int64
DVTY2ST    int64
DVLAST30   int64
Length: 73, dtype: object
```

## CLASSIFICATION

We will work with 2 dataframes:

1. df - a dataframe containing only the primary data.
2. merged\_df - a dataframe containing the primary data and additional context data.

### Model 1 - Classification analysis using a Random Forest Classifier with k-fold cross-validation

Target variable - TS\_011 - Predicted use of cigarettes in the next 12 months

```
from sklearn.model_selection import KFold, cross_val_predict
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, classification_report, confusion_matrix
from sklearn.impute import SimpleImputer

# Define features (X) and target (y)
X = df.drop(columns=['TS_011'])
y = df['TS_011']

# Set up k-Fold Cross-Validation
k = 5 # Number of folds
kf = KFold(n_splits=k, shuffle=True, random_state=42)

# Train a Random Forest classifier
model = RandomForestClassifier(n_estimators=100, random_state=42)

# Generate cross-validated predictions
y_pred = cross_val_predict(model, X, y, cv=kf)

# Calculate and display metrics
print("Accuracy:", accuracy_score(y, y_pred))
print("Precision:", precision_score(y, y_pred, average='weighted'))
print("Recall:", recall_score(y, y_pred, average='weighted'))
print("F1 Score:", f1_score(y, y_pred, average='weighted'))

# Full classification report
print("\nClassification Report:\n", classification_report(y, y_pred))

# Confusion matrix
conf_matrix = confusion_matrix(y, y_pred)
plt.figure(figsize=(6, 4))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=sorted(y.unique()), yticklabels=sorted(y.unique()))
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```




```

Accuracy: 0.819903610383022
Precision: 0.8021377845870976
Recall: 0.819903610383022
F1 Score: 0.8081115078638048

```

Classification Report:					
	precision	recall	f1-score	support	
1	0.68	0.59	0.63	491	
2	0.46	0.21	0.29	864	
3	0.52	0.45	0.48	3952	
4	0.88	0.94	0.91	18347	
accuracy			0.82	23654	
macro avg	0.64	0.55	0.58	23654	
weighted avg	0.80	0.82	0.81	23654	



### Adding Context Data

Read a file containing context data and merge it with the existing CSTADS data based on the PROVID column. Create a new comprehensive dataframe named "merged\_df".

```

# Load CSTADS data
cstads_df = df

# Load context data
context_file = "Context Data 3.csv"
context_df = pd.read_csv(context_file)

# Merge datasets on the 'province_id' column
merged_df = cstads_df.merge(context_df, on="PROVID", how="left")

# Save the merged dataset (optional)
merged_df.to_csv("CSTADS_with_context.csv", index=False)

# Check the first few rows
# print("\n", merged_df.head())
# print(merged_df.dtypes)

```

### Generate EDA Report

Perform analysis on merged dataset to examine correlations

```

from ydata_profiling import ProfileReport

# Create the profile
profile = ProfileReport(merged_df, title="Merged CSTADS EDA Report", explorative=True)

# Save to an HTML file
profile.to_file("merged_cstads_eda_report.html")
files.download("merged_cstads_eda_report.html")

```

 Show hidden output

## Rerun the model with context data

```
# Define features (X) and target (y) using the merged dataset
X = merged_df.drop(columns=['TS_011'])
y = merged_df['TS_011']

# Set up k-Fold Cross-Validation
k = 5 # Number of folds
kf = KFold(n_splits=k, shuffle=True, random_state=42)


# Train a model using cross-validation
model = RandomForestClassifier(n_estimators=100, random_state=42)

# Generate cross-validated predictions
y_pred = cross_val_predict(model, X, y, cv=kf)


# Calculate and display metrics
print("Accuracy:", accuracy_score(y, y_pred))
print("Precision:", precision_score(y, y_pred, average='weighted'))
print("Recall:", recall_score(y, y_pred, average='weighted'))
print("F1 Score:", f1_score(y, y_pred, average='weighted'))

# Full classification report
print("\nClassification Report:\n", classification_report(y, y_pred))

# Confusion matrix
conf_matrix = confusion_matrix(y, y_pred)
plt.figure(figsize=(6, 4))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=sorted(y.unique()), yticklabels=sorted(y.unique()))
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```

 Accuracy: 0.8215946562949185  
Precision: 0.8031068385509972  
Recall: 0.8215946562949185  
F1 Score: 0.8092103986138576

Classification Report:				
	precision	recall	f1-score	support
1	0.67	0.59	0.63	491
2	0.47	0.22	0.30	864
3	0.53	0.44	0.48	3952
4	0.88	0.94	0.91	18347
accuracy			0.82	23654
macro avg	0.64	0.55	0.58	23654
weighted avg	0.80	0.82	0.81	23654



## Model 2 - Classification analysis using Logistic Regression with k-fold cross-validation

Target variable - TS\_011 - Predicted use of cigarettes in the next 12 months

```
from sklearn.model_selection import KFold, cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
import pandas as pd

# Define features (X) and target (y)
X = df.drop(columns=['TS_011'])
y = df['TS_011']

# Feature scaling using StandardScaler
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Set up k-Fold Cross-Validation
k = 5 # Number of folds
kf = KFold(n_splits=k, shuffle=True, random_state=42)

# Initialize Logistic Regression model
model = LogisticRegression(max_iter=5000, random_state=42)

# Generate cross-validated predictions
y_pred = cross_val_predict(model, X_scaled, y, cv=kf)


# Evaluate performance
print("Accuracy:", accuracy_score(y, y_pred))
print("Precision:", precision_score(y, y_pred, average='weighted'))
print("Recall:", recall_score(y, y_pred, average='weighted'))
print("F1 Score:", f1_score(y, y_pred, average='weighted'))

# Full classification report
print("\nClassification Report:\n", classification_report(y, y_pred))

# Confusion matrix
conf_matrix = confusion_matrix(y, y_pred)
plt.figure(figsize=(6, 4))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=sorted(y.unique()), yticklabels=sorted(y.unique()))
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.tight_layout()
plt.show()
```

Accuracy: 0.812336179927285  
 Precision: 0.7863422030689603  
 Recall: 0.812336179927285  
 F1 Score: 0.7949895470476269

Classification Report:					
	precision	recall	f1-score	support	
1	0.66	0.58	0.62	491	
2	0.45	0.25	0.32	864	
3	0.49	0.35	0.40	3952	
4	0.87	0.95	0.91	18347	
accuracy			0.81	23654	
macro avg	0.62	0.53	0.56	23654	
weighted avg	0.79	0.81	0.79	23654	



### Elimination of highly correlated variables

```

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import KFold, cross_val_predict
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, classification_report, confusion_matrix
import numpy as np # Import NumPy

# Define features and target
X = df.drop(columns=['TS_011'])
y = df['TS_011']

# Step: Remove highly correlated features
corr_matrix = X.corr().abs()
upper_triangle = corr_matrix.where(
    np.triu(np.ones(corr_matrix.shape), k=1).astype(bool) # Use np.triu instead of pd.np.triu
)
high_corr_cols = [
    column for column in upper_triangle.columns if any(upper_triangle[column] > 0.85)
]
print("Dropping highly correlated features:", high_corr_cols)
X = X.drop(columns=high_corr_cols)

# Feature scaling
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Set up k-Fold Cross-Validation
kf = KFold(n_splits=5, shuffle=True, random_state=42)

# Logistic Regression model
model = LogisticRegression(max_iter=5000, random_state=42)

# Generate cross-validated predictions
y_pred = cross_val_predict(model, X_scaled, y, cv=kf)
  
```


```
# Evaluate performance
print("Accuracy:", accuracy_score(y, y_pred))
print("Precision:", precision_score(y, y_pred, average='weighted'))
print("Recall:", recall_score(y, y_pred, average='weighted'))
print("F1 Score:", f1_score(y, y_pred, average='weighted'))

print("\nClassification Report:\n", classification_report(y, y_pred))

# Confusion matrix
conf_matrix = confusion_matrix(y, y_pred)
plt.figure(figsize=(6, 4))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=sorted(y.unique()), yticklabels=sorted(y.unique()))
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.tight_layout()
plt.show()
```

↩ Dropping highly correlated features: ['VAP\_050b', 'ELC\_042']  
Accuracy: 0.8127589414052592  
Precision: 0.7868586216674444  
Recall: 0.8127589414052592  
F1 Score: 0.7954458474832734

Classification Report:					
	precision	recall	f1-score	support	
1	0.66	0.58	0.62	491	
2	0.45	0.25	0.32	864	
3	0.49	0.35	0.41	3952	
4	0.87	0.95	0.91	18347	
accuracy			0.81	23654	
macro avg	0.62	0.53	0.56	23654	
weighted avg	0.79	0.81	0.80	23654	



## Rerun the model with context data

```
# Define features (X) and target (y) using merged_df
X = merged_df.drop(columns=['TS_011'])
y = merged_df['TS_011']

# Feature scaling using StandardScaler
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Set up k-Fold Cross-Validation
k = 5 # Number of folds
kf = KFold(n_splits=k, shuffle=True, random_state=42)

# Initialize Logistic Regression model
model = LogisticRegression(max_iter=5000, random_state=42)

# Generate cross-validated predictions
```

```

y_pred = cross_val_predict(model, X_scaled, y, cv=kf)

# Evaluate performance
print("Accuracy:", accuracy_score(y, y_pred))
print("Precision:", precision_score(y, y_pred, average='weighted'))
print("Recall:", recall_score(y, y_pred, average='weighted'))
print("F1 Score:", f1_score(y, y_pred, average='weighted'))

# Full classification report
print("\nClassification Report:\n", classification_report(y, y_pred))

# Confusion matrix
conf_matrix = confusion_matrix(y, y_pred)
plt.figure(figsize=(6, 4))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=sorted(y.unique()), yticklabels=sorted(y.unique()))
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.tight_layout()
plt.show()

```

```

↩ Accuracy: 0.8134353597700178
Precision: 0.7880417528224349
Recall: 0.8134353597700178
F1 Score: 0.7964187192740427

```


```

Classification Report:
              precision    recall  f1-score   support

     1         0.66       0.58       0.62         491
     2         0.46       0.26       0.33         864
     3         0.49       0.35       0.41        3952
     4         0.87       0.95       0.91       18347

 accuracy         0.81       0.81       0.81       23654
 macro avg        0.62       0.53       0.57       23654
 weighted avg     0.79       0.81       0.80       23654

```



## Remove Highly Correlated Features

```

# Define features and target
X = merged_df.drop(columns=['TS_011'])
y = merged_df['TS_011']

# Step: Remove highly correlated features
corr_matrix = X.corr().abs()
upper_triangle = corr_matrix.where(
    np.triu(np.ones(corr_matrix.shape), k=1).astype(bool) # Use np.triu instead of pd.np.triu
)
high_corr_cols = [
    column for column in upper_triangle.columns if any(upper_triangle[column] > 0.85)
]
print("Dropping highly correlated features:", high_corr_cols)
X = X.drop(columns=high_corr_cols)

```

```
# Feature scaling
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Set up k-Fold Cross-Validation
kf = KFold(n_splits=5, shuffle=True, random_state=42)

# Logistic Regression model
model = LogisticRegression(max_iter=5000, random_state=42)

# Generate cross-validated predictions
y_pred = cross_val_predict(model, X_scaled, y, cv=kf)

# Evaluate performance
print("Accuracy:", accuracy_score(y, y_pred))
print("Precision:", precision_score(y, y_pred, average='weighted'))
print("Recall:", recall_score(y, y_pred, average='weighted'))
print("F1 Score:", f1_score(y, y_pred, average='weighted'))

print("\nClassification Report:\n", classification_report(y, y_pred))
```

```
➡ Dropping highly correlated features: ['VAP_050b', 'ELC_042', 'MEANCIGCOST']
Accuracy: 0.8133508074744229
Precision: 0.7879004336508605
Recall: 0.8133508074744229
F1 Score: 0.7963505084125293
```

```
Classification Report:
              precision    recall  f1-score   support

     1         0.66       0.58       0.62         491
     2         0.46       0.25       0.33         864
     3         0.49       0.35       0.41        3952
     4         0.87       0.95       0.91       18347

 accuracy                   0.81       23654
 macro avg                  0.62       23654
 weighted avg               0.79       23654
```

## XGBOOST

```
pip install xgboost
```

```
➡ Requirement already satisfied: xgboost in /usr/local/lib/python3.11/dist-packages (2.1.4)
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (from xgboost) (2.0.2)
Requirement already satisfied: nvidia-nccl-cu12 in /usr/local/lib/python3.11/dist-packages (from xgboost) (2.21.5)
Requirement already satisfied: scipy in /usr/local/lib/python3.11/dist-packages (from xgboost) (1.14.1)
```

```
from xgboost import XGBClassifier
from sklearn.model_selection import KFold, cross_val_predict

# Define features and target
X = df.drop(columns=['TS_011'])
y = df['TS_011']

# Optional: Scale features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Initialize XGBoost classifier
model = XGBClassifier(
    objective='multi:softmax',
    num_class=4,
    use_label_encoder=False,
    eval_metric='mlogloss',
    random_state=42
)

# Set up k-Fold cross-validation
kf = KFold(n_splits=5, shuffle=True, random_state=42)

# Before fitting/predicting, adjust target variable labels to start from 0
y = y - 1

# Generate predictions
```

```
y_pred = cross_val_predict(model, X_scaled, y, cv=kf)
```

```
# Evaluate performance
```

```
print("Accuracy:", accuracy_score(y, y_pred))
print("Precision:", precision_score(y, y_pred, average='weighted'))
print("Recall:", recall_score(y, y_pred, average='weighted'))
print("F1 Score:", f1_score(y, y_pred, average='weighted'))
```

```
print("\nClassification Report:\n", classification_report(y, y_pred))
```

```
# Confusion matrix
```

```
conf_matrix = confusion_matrix(y, y_pred)
plt.figure(figsize=(6, 4))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=sorted(y.unique()), yticklabels=sorted(y.unique()))
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.tight_layout()
plt.show()
```

```
→ /usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [08:56:59] WARNING: /workspace/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.
```

```
warnings.warn(smsg, UserWarning)
/usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [08:57:02] WARNING: /workspace/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.
```

```
warnings.warn(smsg, UserWarning)
/usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [08:57:04] WARNING: /workspace/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.
```

```
warnings.warn(smsg, UserWarning)
/usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [08:57:05] WARNING: /workspace/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.
```


```
warnings.warn(smsg, UserWarning)
/usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [08:57:06] WARNING: /workspace/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.
```

```
warnings.warn(smsg, UserWarning)
Accuracy: 0.8241312251627632
Precision: 0.8145073232860927
Recall: 0.8241312251627632
F1 Score: 0.8184303219374425
```

```
Classification Report:
              precision    recall  f1-score   support

0           0.69         0.57         0.63         491
1           0.42         0.27         0.33         864
2           0.53         0.51         0.52        3952
3           0.90         0.92         0.91       18347

 accuracy          0.82         0.82        23654
 macro avg         0.63         0.57         0.60        23654
 weighted avg      0.81         0.82         0.82        23654
```





## Rerun the model with context data

```
# Define features and target
X = merged_df.drop(columns=['TS_011'])
y = merged_df['TS_011']

# Optional: Scale features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Initialize XGBoost classifier
model = XGBClassifier(
    objective='multi:softmax',
    num_class=4,
    use_label_encoder=False,
    eval_metric='mlogloss',
    random_state=42
)

# Set up k-Fold cross-validation
kf = KFold(n_splits=5, shuffle=True, random_state=42)

# Before fitting/predicting, adjust target variable labels to start from 0
y = y - 1

# Generate predictions
y_pred = cross_val_predict(model, X_scaled, y, cv=kf)

# Evaluate performance
print("Accuracy:", accuracy_score(y, y_pred))
print("Precision:", precision_score(y, y_pred, average='weighted'))
print("Recall:", recall_score(y, y_pred, average='weighted'))
print("F1 Score:", f1_score(y, y_pred, average='weighted'))

print("\nClassification Report:\n", classification_report(y, y_pred))

# Confusion matrix
conf_matrix = confusion_matrix(y, y_pred)
plt.figure(figsize=(6, 4))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=sorted(y.unique()), yticklabels=sorted(y.unique()))
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.tight_layout()
plt.show()
```

```
🔗 /usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [08:57:29] WARNING: /workspace/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.
```

```
warnings.warn(smsg, UserWarning)
/usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [08:57:35] WARNING: /workspace/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.
```

```
warnings.warn(smsg, UserWarning)
/usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [08:57:36] WARNING: /workspace/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.
```

```
warnings.warn(smsg, UserWarning)
/usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [08:57:37] WARNING: /workspace/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.
```


```
warnings.warn(smsg, UserWarning)
/usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [08:57:39] WARNING: /workspace/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.
```

```
warnings.warn(smsg, UserWarning)
Accuracy: 0.8235393590935994
Precision: 0.8132808919553798
Recall: 0.8235393590935994
F1 Score: 0.8173436717077077
```

```
Classification Report:
              precision    recall  f1-score   support

     0       0.67       0.57       0.62         491
     1       0.41       0.24       0.30         864
     2       0.53       0.51       0.52       3952
     3       0.90       0.92       0.91      18347

 accuracy          0.82          23654
 macro avg         0.63          23654
 weighted avg      0.81          23654
```



## Feature Importance

```
model.fit(X, y)


# Get feature importances
importances = model.feature_importances_
feature_names = X.columns
importance_df = pd.DataFrame({'Feature': feature_names, 'Importance': importances})
importance_df = importance_df.sort_values(by='Importance', ascending=False)

# Display top 20 features
plt.figure(figsize=(10, 6))
plt.barh(importance_df['Feature'][:20][::-1], importance_df['Importance'][:20][::-1])
plt.title("Top 20 Most Important Features (XGBoost)")
plt.xlabel("Importance Score")
plt.tight_layout()
plt.show()

# Optional: Print top 20 for inspection
print(importance_df.head(20))
```

```
→ /usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [09:32:23] WARNING: /workspace/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.
```

```
warnings.warn(msg, UserWarning)
```



	Feature	Importance
20	VAP_010	0.257704
71	DVLAST30	0.098375
11	SS_010	0.051852
70	DVTY2ST	0.049963
28	ALC_010	0.020471
34	CAN_010	0.016693
69	DVTY1ST	0.016463
44	PH_020	0.013588
23	VAP_030	0.010355
6	DVORIENT	0.009986
21	CI_010	0.009621
43	PH_010	0.009254
38	PR_100	0.008194
77	CIGTAX	0.008184
59	CA_020	0.008016
54	PH_080	0.007890
9	GH_010	0.007884
60	ELC_041	0.007706
3	DVGENDER	0.007701
17	ELC_026a	0.007687

## CLUSTERING

### Clean and prep the data

```
import pandas as pd
import numpy as np
from google.colab import files

# Load the dataset
df = pd.read_csv("cstads_2122_pumf.csv")

# Define the list of features to keep
features_to_keep = [
    # Smoking-related features
    "SS_010", "SS_020", "TS_011", "SS_030", "SS_040", "WP_040a", "WP_040b", "WP_040c", "WP_040d",
    "WP_040e", "WP_040f", "WP_040g", "SC_010", "CA_011", "TP_001", "TP_016", "TP_046", "TP_056",
    "TP_066", "TP_086",
    # Vaping-related features
```

```

"ELC_026a", "ELC_026b", "ELC_026c", "VAP_010", "CI_010", "VAP_020", "VAP_030", "VAP_040",
"VAP_050a", "VAP_050b", "VAP_060",

# Respondent characteristics
"PROVID", "GRADE", "DVGENDER", "DVURBAN", "DVRES", "DVORIENT", "DVDESCRIBE"
]

# Drop all columns that are NOT in the features_to_keep list
df = df[features_to_keep]

# Remove duplicates if any
df = df.drop_duplicates()

# Identify missing values
df.replace({96: pd.NA, 98: pd.NA, 99: pd.NA, 996: pd.NA, 999: pd.NA}, inplace=True)

# Calculate the percentage of missing values per column
missing_percentage = df.isnull().mean() * 100

# Drop columns with more than 50% missing values
df_cleaned = df.loc[:, missing_percentage <= 50]

# Move cleaned dataframe to df for convenience
df = df_cleaned

# Drop all rows that still contain any missing values
df.dropna(inplace=True)

# Print new dataframe shape to confirm columns are dropped
print("Updated dataframe shape:\n", df.shape)
print("Updated dataframe features:\n", df.dtypes)

```

```

↗ Updated dataframe shape:
(17853, 25)
Updated dataframe features:
SS_010      object
TS_011      object
TP_016      object
TP_046      object
TP_056      object
TP_066      object
TP_086      object
ELC_026a    object
ELC_026b    object
ELC_026c    object
VAP_010     object
CI_010      object
VAP_020     object
VAP_030     object
VAP_040     object
VAP_050a    object
VAP_050b    object
VAP_060     object
PROVID      int64
GRADE       int64
DVGENDER    object
DVURBAN     int64
DVRES       object
DVORIENT    object
DVDESCRIBE  object
dtype: object
<ipython-input-101-f6fa1558978d>:42: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df.dropna(inplace=True)
```

## Identify and drop highly correlated features

```

# Set correlation threshold
threshold = 0.8

# Compute the correlation matrix
corr_matrix = df.corr()

# Find pairs of highly correlated features
high_corr_features = []

```

```

for i in range(len(corr_matrix.columns)):
    for j in range(i):
        if abs(corr_matrix.iloc[i, j]) > threshold: # Check absolute correlation
            high_corr_features.append((corr_matrix.columns[i], corr_matrix.columns[j], corr_matrix.iloc[i, j]))

# Print highly correlated features
if high_corr_features:
    print("Highly Correlated Features (|Correlation| > 0.8):")
    for feature1, feature2, corr_value in high_corr_features:
        print(f"{feature1} ↔ {feature2} | Correlation: {corr_value:.3f}")
else:
    print("No highly correlated features found.")

# Remove one feature from each highly correlated pair identified above
features_to_drop = set()
for feature1, feature2, _ in high_corr_features:
    features_to_drop.add(feature2) # Drop the second feature in the pair

df.drop(columns=features_to_drop, inplace=True)

print(f"Dropped {len(features_to_drop)} feature(s) due to high correlation.")

# Print new dataframe shape to confirm columns are dropped
print("Updated dataframe shape:", df.shape)
print("Updated dataframe features:\n", df.dtypes)

```

```

↔ No highly correlated features found.
Dropped 0 feature(s) due to high correlation.
Updated dataframe shape: (17853, 24)
Updated dataframe features:
SS_010      object
TS_011      object
TP_016      object
TP_046      object
TP_056      object
TP_066      object
TP_086      object
ELC_026a    object
ELC_026b    object
ELC_026c    object
VAP_010     object
CI_010      object
VAP_020     object
VAP_030     object
VAP_040     object
VAP_050b    object
VAP_060     object
PROVID      int64
GRADE       int64
DVGENDER    object
DVURBAN     int64
DVRES       object
DVORIENT    object
DVDESCRIBE  object
dtype: object

```

## Model 1 - DBSCAN Clustering Analysis

```

from sklearn.cluster import DBSCAN
from sklearn.decomposition import PCA
import matplotlib.pyplot as plt

# Apply DBSCAN clustering
dbscan = DBSCAN(eps=3.5, min_samples=5) # Adjust 'eps' based on data distribution
clusters = dbscan.fit_predict(df)

# Add cluster labels to the original dataframe
df['Cluster'] = clusters


# Analyze cluster distribution
print("Cluster distribution:\n", df['Cluster'].value_counts())

# Visualize Clusters using PCA
pca = PCA(n_components=2)
df_pca = pca.fit_transform(df)
plt.scatter(df_pca[:, 0], df_pca[:, 1], c=df['Cluster'], cmap='viridis', alpha=0.5)
plt.xlabel("PCA Component 1")

```

```
plt.ylabel("PCA Component 2")
plt.title("DBSCAN Clustering Visualization")
plt.colorbar(label="Cluster Label")
plt.show()
```

```
Cluster distribution:
Cluster
-1      5001
 3      3585
 0      2187
 1      2056
 2      1288
...
80         5
117        5
89         5
101        5
118        4
Name: count, Length: 120, dtype: int64
```



```
# A dictionary mapping variable names to short descriptions
column_mapping = {
    "SS_010": "Ever Tried Smoking",
    "TS_011": "Likelihood to Smoke",
    "SS_030": "Ever Smoked a Whole Cigarette",
    "SS_040": "Smoked 100+ Cigarettes",
    "WP_040a": "Cigarettes/Day (Day 1)",
    "WP_040b": "Cigarettes/Day (Day 2)",
    "WP_040c": "Cigarettes/Day (Day 3)",
    "WP_040d": "Cigarettes/Day (Day 4)",
    "WP_040e": "Cigarettes/Day (Day 5)",
    "WP_040f": "Cigarettes/Day (Day 6)",
    "WP_040g": "Cigarettes/Day (Day 7)",
    "SC_010": "Tried to Quit Smoking",
    "CA_011": "Usual Cigarette Source",
    "TP_001": "Cigarettes in Last 30 Days",
    "TP_016": "Cigars in Last 30 Days",
    "TP_046": "Smokeless Tobacco in 30 Days",
    "TP_056": "Nicotine Therapy in 30 Days",
    "TP_066": "Hookah Use in 30 Days",
    "TP_086": "Heated Tobacco in 30 Days",
    "ELC_026a": "Vaped Nicotine (30 Days)",
    "ELC_026b": "Vaped No Nicotine (30 Days)",
    "ELC_026c": "Vaped Unknown Substance (30 Days)",
    "VAP_010": "Likelihood to Vape",
    "CI_010": "First Use: Cigarette or Vape?",
    "VAP_020": "Most Used Vape Flavor",
    "VAP_030": "Reason for Trying Vaping",
    "VAP_040": "Reason for Continued Vaping",
    "VAP_050a": "Usual Vape Device Source",
    "VAP_050b": "Usual E-Liquid Source",
    "VAP_060": "Tried to Quit Vaping",
    "SEQID": "Respondent ID",
```

```

    "PROVID": "Province",
    "GRADE": "Grade",
    "DVGENDER": "Gender",
    "DVURBAN": "Urban/Rural School",
    "DVRES": "Years in Canada",
    "DVORIENT": "Sexual Orientation",
    "DVDESCRIBE": "Ethnicity"
}

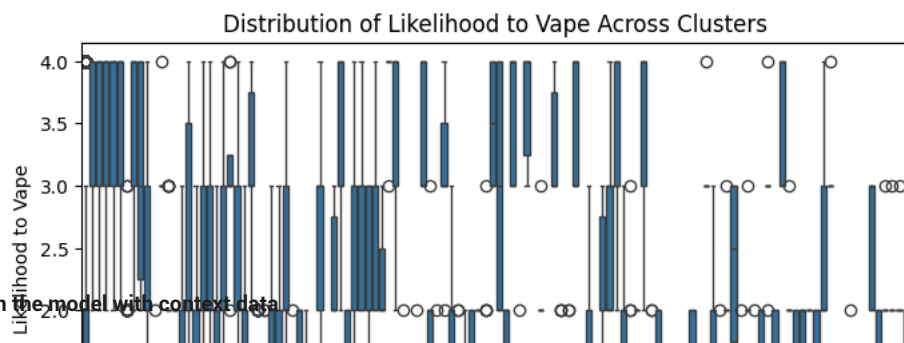
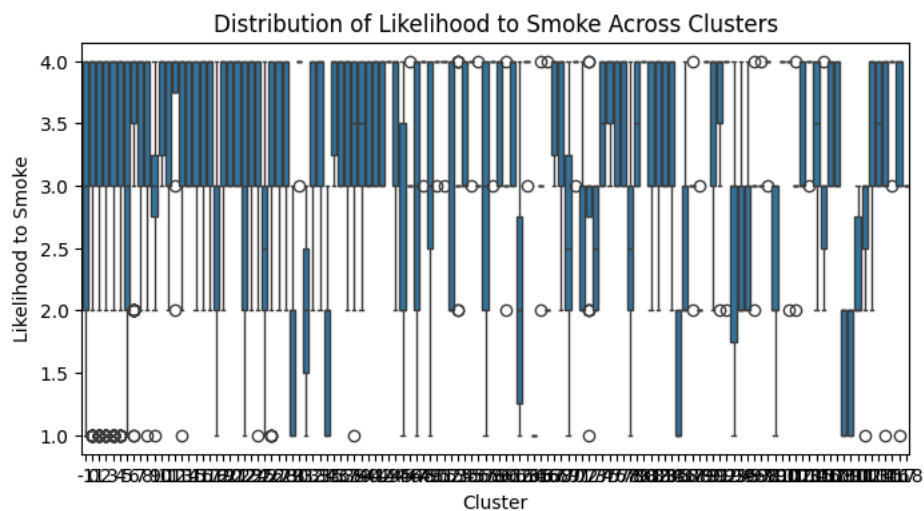

# The variables to visualize
selected_columns = ["Ever Tried Smoking", "Likelihood to Smoke", "Likelihood to Vape", "Grade"]

# Rename columns in the DataFrame
named_df=df.rename(columns=column_mapping)

# Display cluster modes
#cluster_modes = df.groupby('Cluster').agg(pd.Series.mode)
#print(cluster_modes)

# Generate boxplots for the selected columns
for col in selected_columns:
    plt.figure(figsize=(8, 4))
    sns.boxplot(x=named_df['Cluster'], y=named_df[col])
    plt.title(f"Distribution of {col} Across Clusters")
    plt.show()

```



### Rerun the model with context data

```
# Load CSTADS data
cstads_df = df
```

```
# Load context data
context_file = "Context Data.csv"
context_df = pd.read_csv(context_file)
```

```
# Merge datasets on the 'province_id' column
merged_df = cstads_df.merge(context_df, on="PROVID", how="left")
```

```
# Apply DBSCAN clustering
dbscan = DBSCAN(eps=3.5, min_samples=5)
clusters = dbscan.fit_predict(merged_df)
```

```
# Add cluster labels to the original dataframe
merged_df['Cluster'] = clusters
```

```
# Analyze cluster distribution
print("Cluster distribution:\n", merged_df['Cluster'].value_counts())
```




```
# Visualize Clusters using PCA
pca = PCA(n_components=2)
df_pca = pca.fit_transform(merged_df)
plt.scatter(df_pca[:, 0], df_pca[:, 1], c=merged_df['Cluster'], cmap='viridis', alpha=0.5)
plt.xlabel("PCA Component 1")
plt.ylabel("PCA Component 2")
plt.title("DBSCAN Clustering Visualization")
plt.colorbar(label="Cluster Label")
plt.show()
```

Cluster distribution:

```
Cluster
-1      5331
3       3582
0       2169
1       2056
2       1288
...
82         5
87         5
68         5
91         5
80         5
```

Name: count, Length: 94, dtype: int64



## Model 2 - K-Means Clustering Analysis

```
from sklearn.cluster import KMeans
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns


# Define features (X) using the original df dataset
X = df.drop(columns=['TS_011'], errors='ignore') # Exclude the target variable

# Ensure X has no missing values using the most frequent value
imputer = SimpleImputer(strategy='most_frequent')
X = pd.DataFrame(imputer.fit_transform(X), columns=X.columns)

# Feature scaling using StandardScaler
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Determine the optimal number of clusters using the Elbow Method
wcss = [] # Within-cluster sum of squares
for i in range(1, 11):
    kmeans = KMeans(n_clusters=i, random_state=42, n_init=10)
    kmeans.fit(X_scaled)
    wcss.append(kmeans.inertia_)
```

```
# Plot the Elbow Method
plt.figure(figsize=(8, 5))
plt.plot(range(1, 11), wcss, marker='o', linestyle='--')
plt.xlabel('Number of Clusters')
plt.ylabel('WCSS')
plt.title('Elbow Method for Optimal Clusters')
plt.show()
```



```
# Run K-Means
kmeans = KMeans(n_clusters=5, random_state=42)
df['Cluster'] = kmeans.fit_predict(X_scaled)

# Check cluster distribution
print(df['Cluster'].value_counts())

# Most common category per cluster
print(df.groupby('Cluster').apply(lambda x: x.mode().iloc[0]))
```



```
Cluster
0    5612
1    5278
4    3178
2    3083
3     702
Name: count, dtype: int64

0      SS_010  TS_011  TP_016  TP_046  TP_056  TP_066  TP_086  ELC_026a  \
Cluster
0           2       4       5       5       5       5       5       5
1           2       4       5       5       5       5       5       5
2           2       4       5       5       5       5       5       4
3           1       1       4       4       4       4       4       1
4           1       3       5       5       5       5       5       1

0      ELC_026b  ELC_026c  ...  VAP_050b  VAP_060  PROVID  GRADE  DVGENDER  \
Cluster
0           5         5  ...           1         1       24      10         1
1           5         5  ...           1         2       48       8         2
2           5         5  ...           7         3       24      11         1
3           4         5  ...           7         3       48      11         2
4           4         5  ...           7         3       10      11         1

0      DVURBAN  DVRES  DVORIENT  DVDESCRIBE  Cluster
Cluster
0           1       1         2           1         0
1           1       3         2           8         1
2           1       1         2           1         2
3           1       1         2           1         3
4           1       1         2           1         4
```

[5 rows x 25 columns]

```
<ipython-input-129-04196cd6a29f>:9: DeprecationWarning: DataFrameGroupBy.apply operated on the grouping columns. This behavior is deprecated
print(df.groupby('Cluster').apply(lambda x: x.mode().iloc[0]))
```


```

from sklearn.decomposition import PCA

pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)

plt.scatter(X_pca[:, 0], X_pca[:, 1], c=df['Cluster'], cmap='viridis', alpha=0.6)
plt.xlabel("PCA Component 1")
plt.ylabel("PCA Component 2")
plt.title("Cluster Visualization with PCA")
plt.colorbar(label="Cluster")
plt.show()

```



```

# Create a DataFrame showing how much each original feature contributes to PC1 and PC2
pca_loadings = pd.DataFrame(pca.components_, columns=X.columns, index=['PCA1', 'PCA2'])

# Display the loadings
print(pca_loadings.T.sort_values(by='PCA1', ascending=False)) # Sort by PC1 contribution
print(pca_loadings.T.sort_values(by='PCA2', ascending=False)) # Sort by PC2 contribution

```

```

PCA1      PCA2
ELC_026a  0.341374 -0.117446
VAP_010   0.310791 -0.151067
SS_010    0.246812  0.059644
ELC_026b  0.225379  0.050585
ELC_026c  0.221861  0.064475
TP_056    0.210585  0.096700
TP_016    0.201025  0.366911
TP_046    0.169233  0.428861
TP_066    0.157439  0.425117
TP_086    0.148399  0.423811
DVRES     0.085906 -0.036159
DVDESCRIBE 0.084028 -0.053969
PROVID    0.061813 -0.084891
DVGENDER  0.009945 -0.156690
DVORIENT  -0.000461 -0.018479
DVURBAN   -0.024548  0.021373
Cluster   -0.086152 -0.220258
GRADE     -0.098996  0.008158
VAP_020   -0.225282  0.185512
VAP_030   -0.248979  0.071310
VAP_060   -0.257376  0.217523
VAP_050b  -0.264322  0.240404
CI_010    -0.289099  0.068861
VAP_040   -0.311637  0.173684
PCA1      PCA2
TP_046    0.169233  0.428861
TP_066    0.157439  0.425117
TP_086    0.148399  0.423811
TP_016    0.201025  0.366911
VAP_050b  -0.264322  0.240404
VAP_060   -0.257376  0.217523

```